

# BiLSTM

February 20, 2026

## 1 Data Preparation

Pipeline: Raw Text → Data Preprocessing → Model Training → Evaluation → Results

### 1.1 Setup

```
[ ]: import nltk
from nltk.corpus import brown
import numpy as np
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
```

### 1.2 Data Discovery

Load Dataset

```
[ ]: nltk.download('brown')
nltk.download('universal_tagset')
tagged_sents = brown.tagged_sents(tagset='universal')

[nltk_data] Downloading package brown to /root/nltk_data...
[nltk_data]   Unzipping corpora/brown.zip.
[nltk_data] Downloading package universal_tagset to /root/nltk_data...
[nltk_data]   Unzipping taggers/universal_tagset.zip.
```

```
[ ]: print(tagged_sents[0])
```

[('The', 'DET'), ('Fulton', 'NOUN'), ('County', 'NOUN'), ('Grand', 'ADJ'),  
('Jury', 'NOUN'), ('said', 'VERB'), ('Friday', 'NOUN'), ('an', 'DET'),  
('investigation', 'NOUN'), ('of', 'ADP'), ("Atlanta's", 'NOUN'), ('recent',  
'ADJ'), ('primary', 'NOUN'), ('election', 'NOUN'), ('produced', 'VERB'), ('``',  
. .), ('no', 'DET'), ('evidence', 'NOUN'), ("''", '.'), ('that', 'ADP'), ('any',  
'DET'), ('irregularities', 'NOUN'), ('took', 'VERB'), ('place', 'NOUN'), ('.',  
. .)]

Flatten all sentences into a list of (word, tag) tuples

```
[ ]: data = [(word, tag) for sent in tagged_sents for word, tag in sent]
```

Convert data array into pandas dataframe

```
[ ]: df = pd.DataFrame(data, columns=['word', 'tag'])
```

```
[ ]: df.head()
```

```
[ ]:      word  tag
0      The  DET
1  Fulton  NOUN
2  County  NOUN
3   Grand  ADJ
4    Jury  NOUN
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1161192 entries, 0 to 1161191
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  --  
 0   word     1161192 non-null   object 
 1   tag      1161192 non-null   object 
dtypes: object(2)
memory usage: 17.7+ MB
```

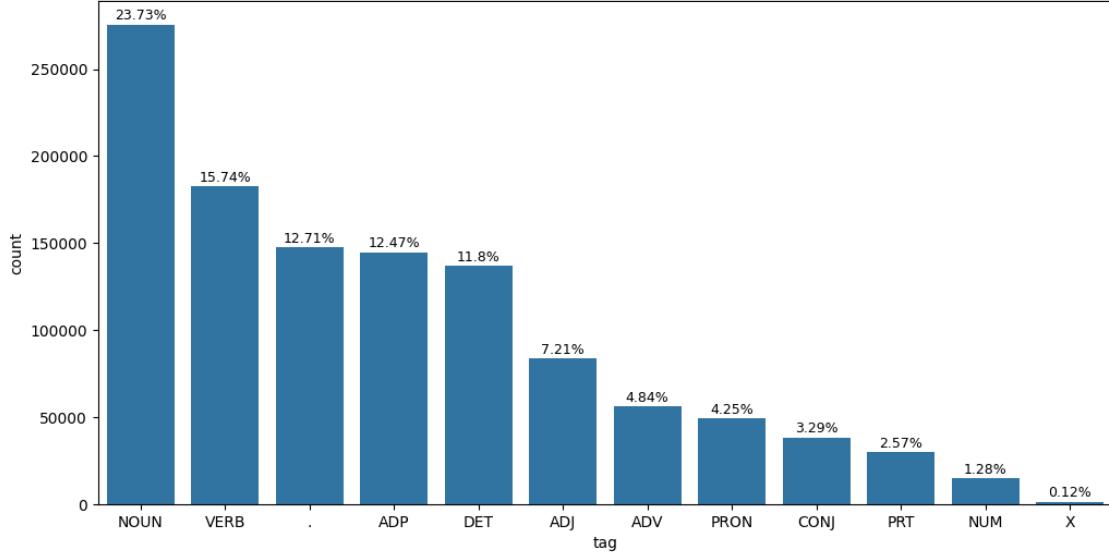
```
[ ]: df['tag'].unique()
```

```
[ ]: array(['DET', 'NOUN', 'ADJ', 'VERB', 'ADP', '.', 'ADV', 'CONJ', 'PRT',
       'PRON', 'NUM', 'X'], dtype=object)
```

```
[ ]: tag_counts = df['tag'].value_counts()
tag_percentages = (tag_counts / len(df) * 100).round(2)

plt.figure(figsize=(12, 6))
ax = sns.countplot(df, x='tag', order=df['tag'].value_counts().index)

for i, (count, pct) in enumerate(zip(tag_counts.values, tag_percentages.
                                         values)):
    ax.text(i, count + 1000, f'{pct}%', ha='center', va='bottom', fontsize=9)
```



## 2 Data Preprocessing

There will be no feature engineering for BiLSTM, but we still have to ensure sentences are whole:

- word: the current word
- tag: the label

### 2.1 Train test split (Split by sentence)

```
[ ]: np.random.seed(42)
n = len(tagged_sents)
indices = np.random.permutation(n)

train_size = int(0.8 * n)
val_size = int(0.1 * n)

train_idx = indices[:train_size]
val_idx = indices[train_size:train_size+val_size]
test_idx = indices[val_size+train_size:]

train_sents = [tagged_sents[i] for i in train_idx]
val_sents = [tagged_sents[i] for i in val_idx]
test_sents = [tagged_sents[i] for i in test_idx]

print(f"\nDataset split:")
print(f"  Train: {len(train_sents)} sentences")
print(f"  Val:   {len(val_sents)} sentences")
```

```
print(f" Test: {len(test_sents):,} sentences")
```

Dataset split:  
Train: 45,872 sentences  
Val: 5,734 sentences  
Test: 5,734 sentences

```
[ ]: print(f"Sample training sentence: {train_sents[0]}")  
print(f"Sample test sentence: {test_sents[0]}")  
print(f"Sample test tags: {val_sents[0]}")
```

```
Sample training sentence: [('Open', 'ADJ'), ('market', 'NOUN'), ('policy', 'NOUN')]  
Sample test sentence: [('Even', 'ADV'), ('the', 'DET'), ('officer', 'NOUN'), ('in', 'ADP'), ('charge', 'NOUN'), ('.', '.'), ('be', 'VERB'), ('it', 'PRON'), ('a', 'DET'), ('captain', 'NOUN'), ('.', '.'), ('for', 'ADP'), ('small', 'ADJ'), ('display', 'NOUN'), ('.', '.'), ('or', 'CONJ'), ('a', 'DET'), ('general', 'NOUN'), ('.', '.'), ('is', 'VERB'), ('restrained', 'VERB'), ('by', 'ADP'), ('monitoring', 'VERB'), ('.', '.')]  
Sample test tags: [('Boys', 'NOUN'), ('will', 'VERB'), ('be', 'VERB'), ('boys', 'NOUN'), ('.', '.'), ('and', 'CONJ'), ('Texans', 'NOUN'), ('will', 'VERB'), ('be', 'VERB'), ('Texans', 'NOUN'), ('.', '.')]
```

### 3 Feature Encoding

Index word and tag vocabularies from sentences. This is to map each word into a number. The vocab is built using the train data.

Args: - **sentences** (list): List of tagged sentences - **min\_freq** (int): Minimum frequency for a word to be included

Returns: - tuple: (word2idx, tag2idx, idx2word, idx2tag)

```
[ ]: from collections import Counter
```

```
[ ]: def build_vocabularies(sentences, min_freq=2):  
  
    # Count word frequencies  
    word_counter = Counter()  
    tag_counter = Counter()  
  
    for sent in sentences:  
        for word, tag in sent:  
            word_counter[word.lower()] += 1  
            tag_counter[tag] += 1  
  
    # Build word vocabulary (only words with freq >= min_freq)
```

```

    words = ['<PAD>', '<UNK>'] + [w for w, c in word_counter.items() if c >= min_freq]
    word2idx = {w: i for i, w in enumerate(words)}
    idx2word = {i: w for w, i in word2idx.items()}

    # Build tag vocabulary
    tags = ['<PAD>'] + sorted(tag_counter.keys())
    tag2idx = {t: i for i, t in enumerate(tags)}
    idx2tag = {i: t for t, i in tag2idx.items()}

    return word2idx, tag2idx, idx2word, idx2tag

```

```
[ ]: word2idx, tag2idx, idx2word, idx2tag = build_vocabularies(train_sents)
print(f"Vocabulary size: {len(word2idx)}")
print(f"Number of POS tags: {len(tag2idx)}")
print(f"POS tags: {list(tag2idx.keys())}")
```

Vocabulary size: 24716  
Number of POS tags: 13  
POS tags: ['<PAD>', '.', 'ADJ', 'ADP', 'ADV', 'CONJ', 'DET', 'NOUN', 'NUM',  
'PRON', 'PRT', 'VERB', 'X']

### 3.1 Encoding Sentences

Convert sentences to numerical indices and pad them. This function uses the vocab dictionary we built earlier to encode the tagged sentences. Words that do not exist in the vocab dictionary appear as ‘UNK’, while padding is to ensure all sequences in a batch have the same shape. (The NN will not be able to process batches with irregular shape)

Args: - `sentences` (list): List of tagged sentences - `word2idx` (dict): Word to index mapping - `tag2idx` (dict): Tag to index mapping - `max_len` (int): Maximum sequence length (computed if None)

Returns: - tuple: (`word_indices`, `tag_indices`, `sequence_lengths`, `max_len`)

```
[ ]: def encode_sentences(sentences, word2idx, tag2idx, max_len=None):

    # Find max length if not provided
    if max_len is None:
        max_len = max(len(sent) for sent in sentences)

    word_indices = []
    tag_indices = []
    seq_lengths = []

    for sent in sentences:
        # Get original length
        seq_len = len(sent)
        seq_lengths.append(seq_len)

        word_indices.append([word2idx.get(w, word2idx['UNK']) for w in sent])
        tag_indices.append([tag2idx.get(t, tag2idx['UNK']) for t in sent])
```

```

# Convert words and tags to indices
words = [word2idx.get(word.lower(), word2idx['<UNK>']) for word, _ in sent]
tags = [tag2idx[tag] for _, tag in sent]

# Pad sequences
if len(words) < max_len:
    words += [word2idx['<PAD>']] * (max_len - len(words))
    tags += [tag2idx['<PAD>']] * (max_len - len(tags))
else:
    # Truncate if needed
    words = words[:max_len]
    tags = tags[:max_len]
    seq_lengths[-1] = max_len

word_indices.append(words)
tag_indices.append(tags)

return np.array(word_indices), np.array(tag_indices), np.
array(seq_lengths), max_len

```

Encode all datasets

```
[ ]: train_X, train_y, train_lengths, max_len = encode_sentences(train_sents, word2idx, tag2idx)
val_X, val_y, val_lengths, _ = encode_sentences(val_sents, word2idx, tag2idx, max_len)
test_X, test_y, test_lengths, _ = encode_sentences(test_sents, word2idx, tag2idx, max_len)
```

```
[ ]: print(f"Max sequence length: {max_len}")
print()
print(f"Training data shape: {train_X.shape}")
print(f"Training labels shape: {train_y.shape}")
print()
print(f"Validation data shape: {val_X.shape}")
print(f"Validation labels shape: {val_y.shape}")
print()
print(f"Test data shape: {test_X.shape}")
print(f"Test labels shape: {test_y.shape}")
```

Max sequence length: 172

Training data shape: (45872, 172)  
 Training labels shape: (45872, 172)

Validation data shape: (5734, 172)

```
Validation labels shape: (5734, 172)
```

```
Test data shape: (5734, 172)
```

```
Test labels shape: (5734, 172)
```

### 3.2 Create dataloaders

Create a PyTorch DataLoader from numpy arrays.

Args: - X (np.array): Input sequences - y (np.array): Target sequences - lengths (np.array): Actual sequence lengths - batch\_size (int): Batch size = 64 (common default) - shuffle (bool): Whether to shuffle data

Returns: - DataLoader: PyTorch DataLoader object

```
[ ]: import torch
      from torch.utils.data import DataLoader, TensorDataset

[ ]: def create_dataloader(X, y, lengths, batch_size=64, shuffle=True):

    # Convert to PyTorch tensors
    X_tensor = torch.LongTensor(X)
    y_tensor = torch.LongTensor(y)
    lengths_tensor = torch.LongTensor(lengths)

    # Create dataset
    dataset = TensorDataset(X_tensor, y_tensor, lengths_tensor)

    # Create dataloader
    dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=shuffle)

    return dataloader

[ ]: train_loader = create_dataloader(train_X, train_y, train_lengths, shuffle=True)
      val_loader = create_dataloader(val_X, val_y, val_lengths, shuffle=False)
      test_loader = create_dataloader(test_X, test_y, test_lengths, shuffle=False)

[ ]: print(f"Number of training batches: {len(train_loader)}")
      print(f"Number of validation batches: {len(val_loader)}")
      print(f"Number of test batches: {len(test_loader)}")
```

```
Number of training batches: 717
```

```
Number of validation batches: 90
```

```
Number of test batches: 90
```

## 4 Build BiLSTM Model

The model has bidirectional context (looks both forward and backward in the sentence). It is also designed to have long-range dependencies (captures relationships across distant words).

## 4.1 Model Architecture:

1. Embedding Layer
  - Transforms encoded words into dense vectors
  - Creates a lookup table where each word gets a unique vector, which are trained during backpropagation
  - Input: one-hot encoded words (word indices)
  - Output: 100-dimensional vectors for each word
2. BiLSTM Layers (2 layers)
  - Captures sequential patterns and bidirectional context (the brain layer)
  - Addresses lexical ambiguity in the English language
  - Input: Word embeddings
  - Output: 256-dimensional hidden states for each word ( $128 \times 2$  bidirectional = 256)
3. Dropout Layer (Only Active for training)
  - Prevents overfitting by randomly dropping neurons during training
  - Helps the model generalize better by learning patterns and not noise
4. Fully Connected Layer
  - Transforms BiLSTM hidden states into tag scores
  - Maps LSTM features to tags
  - Input: Hidden states
  - Output: Tag scores (for each word -> 13 scores)
5. Output Layer (Softmax)
  - Converts scores into probabilities
  - Input: Tag scores
  - Output: All tag probabilities for each word

Args: - `vocab_size` (int): Size of vocabulary - `embedding_dim` (int): Dimension of word embeddings - `hidden_dim` (int): Dimension of LSTM hidden state - `output_size` (int): Number of output tags - `num_layers` (int): Number of LSTM layers - `dropout` (float): Dropout probability

Returns: - `nn.Module`: PyTorch model

```
[ ]: import torch.nn as nn

[ ]: def create_bilstm_model(vocab_size, embedding_dim, hidden_dim, output_size,
                           num_layers=2, dropout=0.3):

    model = nn.Sequential()

    class BiLSTM(nn.Module):
        def __init__(self):
            super(BiLSTM, self).__init__()

            # Embedding layer
            self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)

            # BiLSTM layer
            self.lstm = nn.LSTM(
```

```

        embedding_dim,
        hidden_dim,
        num_layers=num_layers,
        bidirectional=True,
        dropout=dropout if num_layers > 1 else 0,
        batch_first=True
    )

    # Dropout layer
    self.dropout = nn.Dropout(dropout)

    # Output layer (hidden_dim * 2 because bidirectional)
    self.fc = nn.Linear(hidden_dim * 2, output_size)

    def forward(self, x, lengths):
        # x shape: (batch_size, seq_len)
        # lengths shape: (batch_size)
        batch_size, seq_len = x.size()
        # Embedding: (batch_size, seq_len, embedding_dim)
        embedded = self.embedding(x)

        # Pack padded sequence for efficient processing
        packed = nn.utils.rnn.pack_padded_sequence(
            embedded, lengths.cpu(), batch_first=True, enforce_sorted=False
        )

        # BiLSTM: (batch_size, seq_len, hidden_dim * 2)
        lstm_out, _ = self.lstm(packed)

        # Unpack sequence
        lstm_out, _ = nn.utils.rnn.pad_packed_sequence(
            lstm_out,
            batch_first=True,
            total_length=seq_len
        )

        # Apply dropout
        lstm_out = self.dropout(lstm_out)

        # Output projection: (batch_size, seq_len, output_size)
        output = self.fc(lstm_out)

    return output

    return BiLSTM()

# Model hyperparameters
EMBEDDING_DIM = 100

```

```
HIDDEN_DIM = 128
NUM_LAYERS = 2
DROPOUT = 0.3
```

```
[ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = create_bilstm_model(
    vocab_size=len(word2idx),
    embedding_dim=EMBEDDING_DIM,
    hidden_dim=HIDDEN_DIM,
    output_size=len(tag2idx),
    num_layers=NUM_LAYERS,
    dropout=DROPOUT
).to(device)
```

```
[ ]: print(f"Device: {device}")
print(f"Model Architecture:\n{model}")
print(f"\nTotal Parameters: {sum(p.numel() for p in model.parameters())}")
```

```
Device: cpu
Model Architecture:
BiLSTM(
    (embedding): Embedding(24716, 100, padding_idx=0)
    (lstm): LSTM(100, 128, num_layers=2, batch_first=True, dropout=0.3,
bidirectional=True)
    (dropout): Dropout(p=0.3, inplace=False)
    (fc): Linear(in_features=256, out_features=13, bias=True)
)
```

```
Total Parameters: 3105725
```

## 5 Training and Evaluation

### 5.1 Loss function and optimizer setup

- **criterion:** Cross Entropy Loss function with ignore\_index for padding (ignore ‘PAD’ tokens)
- **optimizer:** Adam optimizer with 0.001 learning rate

```
[ ]: import torch.optim as optim
```

```
[ ]: criterion = nn.CrossEntropyLoss(ignore_index=tag2idx['<PAD>'])
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

Train the model for one epoch

```
[ ]: def train_epoch(model, dataloader, criterion, optimizer, device):
    model.train()
    total_loss = 0
```

```

total_correct = 0
total_tokens = 0

for batch_idx, (inputs, targets, lengths) in enumerate(dataloader):
    # Move to device
    inputs = inputs.to(device)
    targets = targets.to(device)
    lengths = lengths.to(device)

    # Zero gradients
    optimizer.zero_grad()

    # Forward pass
    outputs = model(inputs, lengths)  # (batch_size, seq_len, num_tags)

    # Reshape for loss calculation
    outputs_flat = outputs.view(-1, outputs.shape[-1])  # (batch_size * ↵seq_len, num_tags)
    targets_flat = targets.view(-1)  # (batch_size * seq_len)

    # Calculate loss
    loss = criterion(outputs_flat, targets_flat)

    # Backward pass
    loss.backward()

    # Clip gradients to prevent exploding gradients
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=5.0)

    # Update weights
    optimizer.step()

    # Calculate accuracy (excluding padding)
    predictions = outputs.argmax(dim=-1)  # (batch_size, seq_len)
    mask = targets != tag2idx['<PAD>']  # Create mask for non-padding tokens

    correct = (predictions == targets) & mask
    total_correct += correct.sum().item()
    total_tokens += mask.sum().item()
    total_loss += loss.item()

    # Print progress every 100 batches
    if (batch_idx + 1) % 100 == 0:
        print(f"  Batch {batch_idx + 1}/{len(dataloader)}, Loss: {loss.item():.4f}")

avg_loss = total_loss / len(dataloader)

```

```

accuracy = total_correct / total_tokens

return avg_loss, accuracy

```

[ ]: from sklearn.metrics import accuracy\_score, classification\_report, confusion\_matrix

Evaluate the model on a dataset

```

[ ]: def evaluate(model, dataloader, criterion, device):

    model.eval()
    total_loss = 0
    all_predictions = []
    all_targets = []

    with torch.no_grad():
        for inputs, targets, lengths in dataloader:
            # Move to device
            inputs = inputs.to(device)
            targets = targets.to(device)
            lengths = lengths.to(device)

            # Forward pass
            outputs = model(inputs, lengths)

            # Calculate loss
            outputs_flat = outputs.view(-1, outputs.shape[-1])
            targets_flat = targets.view(-1)
            loss = criterion(outputs_flat, targets_flat)
            total_loss += loss.item()

            # Get predictions
            predictions = outputs.argmax(dim=-1)

            # Store predictions and targets (excluding padding)
            for i in range(len(inputs)):
                seq_len = lengths[i].item()
                pred_seq = predictions[i, :seq_len].cpu().numpy()
                target_seq = targets[i, :seq_len].cpu().numpy()

                all_predictions.extend(pred_seq)
                all_targets.extend(target_seq)

    avg_loss = total_loss / len(dataloader)
    accuracy = accuracy_score(all_targets, all_predictions)

    return avg_loss, accuracy, all_predictions, all_targets

```

Train model on specified epochs

```
[ ]: def train_model(model, train_loader, val_loader, criterion, optimizer, device, num_epochs=10):

    history = {
        'train_loss': [],
        'train_acc': [],
        'val_loss': [],
        'val_acc': []
    }

    best_val_acc = 0.0

    for epoch in range(num_epochs):
        print(f"\nEpoch {epoch + 1}/{num_epochs}")
        print("-" * 50)

        # Train
        train_loss, train_acc = train_epoch(model, train_loader, criterion, optimizer, device)
        print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}")

        # Validate
        val_loss, val_acc, _, _ = evaluate(model, val_loader, criterion, device)
        print(f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}")

        # Save history
        history['train_loss'].append(train_loss)
        history['train_acc'].append(train_acc)
        history['val_loss'].append(val_loss)
        history['val_acc'].append(val_acc)

        # Save best model
        if val_acc > best_val_acc:
            best_val_acc = val_acc
            print(f" → New best validation accuracy: {best_val_acc:.4f}")

    return history
```

Train the model

```
[ ]: import time

[ ]: training_start = time.perf_counter()
history = train_model(model, train_loader, val_loader, criterion, optimizer, device)
training_end = time.perf_counter()
```

Epoch 1/10

---

```
Batch 100/717, Loss: 0.6550
Batch 200/717, Loss: 0.4209
Batch 300/717, Loss: 0.3611
Batch 400/717, Loss: 0.2586
Batch 500/717, Loss: 0.2778
Batch 600/717, Loss: 0.2416
Batch 700/717, Loss: 0.2398
Train Loss: 0.4641, Train Acc: 0.8471
Val Loss: 0.1891, Val Acc: 0.9371
→ New best validation accuracy: 0.9371
```

Epoch 2/10

---

```
Batch 100/717, Loss: 0.1581
Batch 200/717, Loss: 0.1802
Batch 300/717, Loss: 0.1430
Batch 400/717, Loss: 0.1625
Batch 500/717, Loss: 0.1554
Batch 600/717, Loss: 0.1468
Batch 700/717, Loss: 0.1299
Train Loss: 0.1690, Train Acc: 0.9447
Val Loss: 0.1229, Val Acc: 0.9595
→ New best validation accuracy: 0.9595
```

Epoch 3/10

---

```
Batch 100/717, Loss: 0.1065
Batch 200/717, Loss: 0.1078
Batch 300/717, Loss: 0.1089
Batch 400/717, Loss: 0.0992
Batch 500/717, Loss: 0.1244
Batch 600/717, Loss: 0.1036
Batch 700/717, Loss: 0.0859
Train Loss: 0.1141, Train Acc: 0.9629
Val Loss: 0.0987, Val Acc: 0.9674
→ New best validation accuracy: 0.9674
```

Epoch 4/10

---

```
Batch 100/717, Loss: 0.0786
Batch 200/717, Loss: 0.0973
Batch 300/717, Loss: 0.0612
Batch 400/717, Loss: 0.0715
Batch 500/717, Loss: 0.0616
Batch 600/717, Loss: 0.1010
```

Batch 700/717, Loss: 0.0799  
Train Loss: 0.0869, Train Acc: 0.9715  
Val Loss: 0.0877, Val Acc: 0.9705  
→ New best validation accuracy: 0.9705

Epoch 5/10

---

Batch 100/717, Loss: 0.0808  
Batch 200/717, Loss: 0.0659  
Batch 300/717, Loss: 0.0576  
Batch 400/717, Loss: 0.0644  
Batch 500/717, Loss: 0.0584  
Batch 600/717, Loss: 0.0848  
Batch 700/717, Loss: 0.1022  
Train Loss: 0.0700, Train Acc: 0.9771  
Val Loss: 0.0813, Val Acc: 0.9731  
→ New best validation accuracy: 0.9731

Epoch 6/10

---

Batch 100/717, Loss: 0.0641  
Batch 200/717, Loss: 0.0344  
Batch 300/717, Loss: 0.0387  
Batch 400/717, Loss: 0.0520  
Batch 500/717, Loss: 0.0528  
Batch 600/717, Loss: 0.0436  
Batch 700/717, Loss: 0.0530  
Train Loss: 0.0588, Train Acc: 0.9805  
Val Loss: 0.0784, Val Acc: 0.9743  
→ New best validation accuracy: 0.9743

Epoch 7/10

---

Batch 100/717, Loss: 0.0396  
Batch 200/717, Loss: 0.0444  
Batch 300/717, Loss: 0.0368  
Batch 400/717, Loss: 0.0346  
Batch 500/717, Loss: 0.0538  
Batch 600/717, Loss: 0.0695  
Batch 700/717, Loss: 0.0512  
Train Loss: 0.0496, Train Acc: 0.9834  
Val Loss: 0.0794, Val Acc: 0.9743

Epoch 8/10

---

Batch 100/717, Loss: 0.0352  
Batch 200/717, Loss: 0.0271  
Batch 300/717, Loss: 0.0485

```
Batch 400/717, Loss: 0.0377
Batch 500/717, Loss: 0.0408
Batch 600/717, Loss: 0.0396
Batch 700/717, Loss: 0.0374
Train Loss: 0.0428, Train Acc: 0.9856
Val Loss: 0.0793, Val Acc: 0.9753
→ New best validation accuracy: 0.9753
```

Epoch 9/10

---

```
Batch 100/717, Loss: 0.0385
Batch 200/717, Loss: 0.0288
Batch 300/717, Loss: 0.0288
Batch 400/717, Loss: 0.0345
Batch 500/717, Loss: 0.0406
Batch 600/717, Loss: 0.0325
Batch 700/717, Loss: 0.0333
Train Loss: 0.0368, Train Acc: 0.9875
Val Loss: 0.0814, Val Acc: 0.9752
```

Epoch 10/10

---

```
Batch 100/717, Loss: 0.0295
Batch 200/717, Loss: 0.0303
Batch 300/717, Loss: 0.0242
Batch 400/717, Loss: 0.0347
Batch 500/717, Loss: 0.0366
Batch 600/717, Loss: 0.0422
Batch 700/717, Loss: 0.0388
Train Loss: 0.0320, Train Acc: 0.9891
Val Loss: 0.0852, Val Acc: 0.9748
```

Classification Report

```
[ ]: inference_start = time.perf_counter()
test_loss, test_acc, test_predictions, test_targets = evaluate(model,
    test_loader, criterion, device)
inference_end = time.perf_counter()

print(f"Training time: {training_end - training_start:.4f}s")
print(f"Inference time: {inference_end - inference_start:.4f}s")
```

Training time: 4215.5727s  
Inference time: 6.7825s

```
[ ]: print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")
```

Test Loss: 0.0908

Test Accuracy: 0.9737

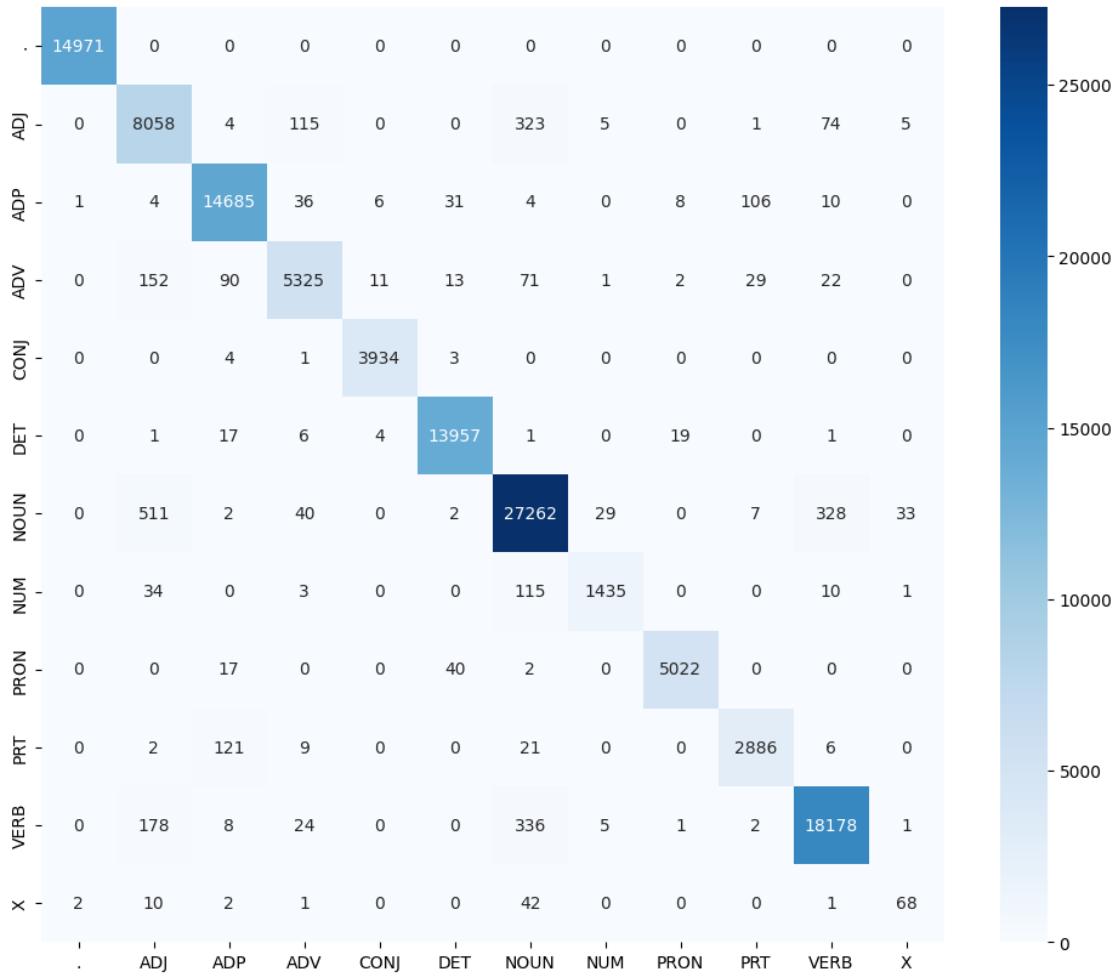
```
[ ]: tag_names = [tag for tag in sorted(tag2idx.keys()) if tag != '<PAD>'] # exclude ↴padding
      print(classification_report(test_targets, test_predictions,
                                    target_names=tag_names,
                                    labels=[tag2idx[tag] for tag in tag_names],
                                    zero_division=0))
```

	precision	recall	f1-score	support
.	1.00	1.00	1.00	14971
ADJ	0.90	0.94	0.92	8585
ADP	0.98	0.99	0.98	14891
ADV	0.96	0.93	0.94	5716
CONJ	0.99	1.00	1.00	3942
DET	0.99	1.00	1.00	14006
NOUN	0.97	0.97	0.97	28214
NUM	0.97	0.90	0.93	1598
PRON	0.99	0.99	0.99	5081
PRT	0.95	0.95	0.95	3045
VERB	0.98	0.97	0.97	18733
X	0.63	0.54	0.58	126
accuracy			0.97	118908
macro avg	0.94	0.93	0.94	118908
weighted avg	0.97	0.97	0.97	118908

Confusion Matrix

```
[ ]: import seaborn as sns
      import matplotlib.pyplot as plt

[ ]: cm = confusion_matrix(test_targets, test_predictions, labels=[tag2idx[tag] for ↴tag in tag_names])
      plt.figure(figsize=(12, 10))
      sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=tag_names, ↴
      yticklabels=tag_names)
      plt.show()
```



## Training and validation metrics

```
[ ]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Plot loss
axes[0].plot(history['train_loss'], label='Train Loss', marker='o')
axes[0].plot(history['val_loss'], label='Val Loss', marker='s')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Loss')
axes[0].set_title('Training and Validation Loss')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Plot accuracy
axes[1].plot(history['train_acc'], label='Train Accuracy', marker='o')
axes[1].plot(history['val_acc'], label='Val Accuracy', marker='s')
axes[1].set_xlabel('Epoch')
```

```

axes[1].set_ylabel('Accuracy')
axes[1].set_title('Training and Validation Accuracy')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```

