

CRF

February 20, 2026

1 Conditional Random Field (CRF) for Part-of-Speech Tagging

Pipeline: Data Preparation → Feature Engineering → Training Phase → Prediction Phase → Evaluation

Key Features: - 25+ hand-crafted linguistic features per word - L-BFGS optimization with L1/L2 regularization - Viterbi decoding for global inference - 97-98% accuracy on Brown Corpus

2 Implementation

2.1 Setup and Installation

```
[ ]: # Install required package
!pip install sklearn-crfsuite

Collecting sklearn-crfsuite
  Using cached sklearn_crfsuite-0.5.0-py2.py3-none-any.whl.metadata (4.9 kB)
Requirement already satisfied: python-crfsuite>=0.9.7 in
/usr/local/lib/python3.12/dist-packages (from sklearn-crfsuite) (0.9.12)
Requirement already satisfied: scikit-learn>=0.24.0 in
/usr/local/lib/python3.12/dist-packages (from sklearn-crfsuite) (1.6.1)
Requirement already satisfied: tabulate>=0.4.2 in
/usr/local/lib/python3.12/dist-packages (from sklearn-crfsuite) (0.9.0)
Requirement already satisfied: tqdm>=2.0 in /usr/local/lib/python3.12/dist-
packages (from sklearn-crfsuite) (4.67.3)
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.12/dist-
packages (from scikit-learn>=0.24.0->sklearn-crfsuite) (2.0.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-
packages (from scikit-learn>=0.24.0->sklearn-crfsuite) (1.16.3)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-
packages (from scikit-learn>=0.24.0->sklearn-crfsuite) (1.5.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.24.0->sklearn-
crfsuite) (3.6.0)
Using cached sklearn_crfsuite-0.5.0-py2.py3-none-any.whl (10 kB)
Installing collected packages: sklearn-crfsuite
Successfully installed sklearn-crfsuite-0.5.0
```

```
[ ]: import nltk
from nltk.corpus import brown
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, precision_recall_fscore_support
import sklearn_crfsuite
from sklearn_crfsuite import metrics
import time
from collections import Counter

print("=="*80)
print("CONDITIONAL RANDOM FIELD (CRF) FOR PART-OF-SPEECH TAGGING")
print("=="*80)
print("Model Type: Discriminative Probabilistic Sequence Model")
print("Optimization: L-BFGS")
print("Features: Hand-crafted linguistic features (~25-30 per word)")
print("=="*80)
```

=====

CONDITIONAL RANDOM FIELD (CRF) FOR PART-OF-SPEECH TAGGING

=====

Model Type: Discriminative Probabilistic Sequence Model
 Optimization: L-BFGS
 Features: Hand-crafted linguistic features (~25-30 per word)

=====

2.2 Step 1: Data Preparation

```
[ ]: print("\n" + "=="*80)
print("STEP 1: DATA PREPARATION")
print("=="*80)

nltk.download('brown')
nltk.download('universal_tagset')
tagged_sents = brown.tagged_sents(tagset='universal')

print(f" Dataset: Brown Corpus")
print(f" Total sentences: {len(tagged_sents)}")
print(f" Tagset: Universal (12 tags)")
print(f" Example sentence: {tagged_sents[0][:5]}...")
print("=="*80)
```

STEP 1: DATA PREPARATION

```
[nltk_data] Downloading package brown to /root/nltk_data...
[nltk_data]   Unzipping corpora/brown.zip.
[nltk_data] Downloading package universal_tagset to /root/nltk_data...
[nltk_data]   Unzipping taggers/universal_tagset.zip.

Dataset: Brown Corpus
Total sentences: 57,340
Tagset: Universal (12 tags)
Example sentence: [('The', 'DET'), ('Fulton', 'NOUN'), ('County', 'NOUN'),
('Grand', 'ADJ'), ('Jury', 'NOUN')]...
```

2.3 Step 2: Data Exploration & Analysis

```
[ ]: print("\n" + "="*80)
print("STEP 2: DATA EXPLORATION & ANALYSIS")
print("="*80)

# Flatten all sentences into a list of (word, tag) tuples
data = [(word, tag) for sent in tagged_sents for word, tag in sent]
df = pd.DataFrame(data, columns=['word', 'tag'])

print(f" Total tokens: {len(df)}")
print(f" Unique words: {df['word'].nunique()}")
print(f" Unique POS tags: {len(df['tag'].unique())}")
print(f" POS tags: {', '.join(sorted(df['tag'].unique()))}")

# Tag distribution
tag_counts = df['tag'].value_counts()
tag_percentages = (tag_counts / len(df) * 100).round(2)

print("\nTag Distribution:")
for tag, count, pct in zip(tag_counts.index[:5], tag_counts.values[:5], tag_percentages.values[:5]):
    print(f" {tag:6s}: {count:7,} ({pct:5.2f}%)")
print(f" ... (7 more tags)")

plt.figure(figsize=(14, 6))
ax = sns.countplot(data=df, x='tag', order=df['tag'].value_counts().index,
                    palette='viridis')
for i, (count, pct) in enumerate(zip(tag_counts.values, tag_percentages.values)):
    ax.text(i, count + 1000, f'{pct}%', ha='center', va='bottom', fontsize=10,
            fontweight='bold')
```

```

plt.title('POS Tag Distribution in Brown Corpus', fontsize=16, u
    ↪fontweight='bold', pad=20)
plt.xlabel('POS Tag', fontsize=12, fontweight='bold')
plt.ylabel('Count', fontsize=12, fontweight='bold')
plt.xticks(rotation=45)
plt.grid(axis='y', alpha=0.3, linestyle='---')
plt.tight_layout()
plt.show()

print("="*80)

```

=====

STEP 2: DATA EXPLORATION & ANALYSIS

=====

```

Total tokens: 1,161,192
Unique words: 56,057
Unique POS tags: 12
POS tags: ., ADJ, ADP, ADV, CONJ, DET, NOUN, NUM, PRON, PRT, VERB, X

```

Tag Distribution:

```

NOUN   : 275,558 (23.73%)
VERB   : 182,750 (15.74%)
.      : 147,565 (12.71%)
ADP    : 144,766 (12.47%)
DET    : 137,019 (11.80%)
... (7 more tags)

```

/tmp/ipython-input-928917345.py:24: FutureWarning:

```

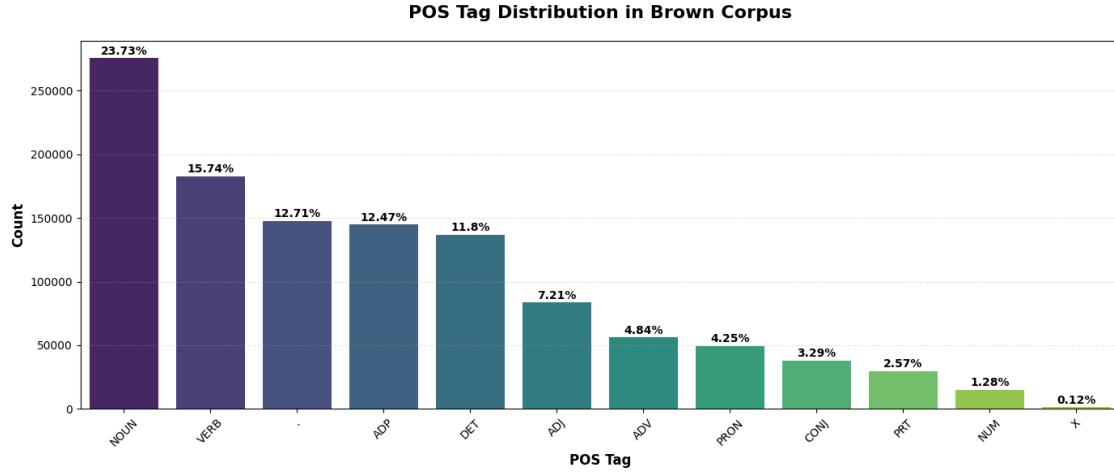
Passing `palette` without assigning `hue` is deprecated and will be removed in
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same
effect.

```

```

ax = sns.countplot(data=df, x='tag', order=df['tag'].value_counts().index,
palette='viridis')

```



2.4 Step 3: Feature Engineering (The Brain of CRF)

This is the core of CRF - rich feature engineering replaces learned embeddings in neural networks.
 #### Feature Engineering Layer (Feature Extraction) **The “brain” of CRF** - Extracts ~25+ linguistic features per word

Unlike BiLSTM’s learned embeddings, these are hand-crafted features:

a) WORD-LEVEL FEATURES:

- Lowercased word form (identity feature)
- Word shape (capitalization patterns: isupper, istitle, isdigit)

b) MORPHOLOGICAL FEATURES (Prefix & Suffix):

- **Prefixes:** word[:1], word[:2], word[:3]
- **Suffixes:** word[-1], word[-2:], word[-3:]
- Helps identify word classes (e.g., “-ing” → VERB, “-ly” → ADV)

c) CHARACTER-LEVEL FEATURES:

- Word length
- Contains hyphen (compound words)
- Contains digit (numerical expressions)
- Contains uppercase letters
- Is alphabetic/alphanumeric
- Contains apostrophe (contractions/possessives)

d) CONTEXTUAL FEATURES (Bidirectional Context):

- Previous word (-1): lowercased, shape, length
- Next word (+1): lowercased, shape, length

- Beginning of Sentence (BOS) marker
- End of Sentence (EOS) marker

Input: List of (word, tag) tuples for a sentence

Output: Feature dictionary for each word position

Feature Parameters: - Prefix lengths: [1, 2, 3] characters - Suffix lengths: [1, 2, 3] characters - Context window: ± 1 word - Total features per word: ~25-30 (varies by position and word properties)

```
[ ]: def word2features(sent, i):
    """
    Extract comprehensive linguistic features for word at position i.

    Args:
        sent: List of (word, tag) tuples representing a sentence
        i: Position of current word (0-indexed)

    Returns:
        dict: Feature dictionary with ~25-30 features
    """
    word = sent[i][0]

    features = {
        # BIAS TERM (always 1.0, allows model to learn tag priors)
        'bias': 1.0,

        # === WORD IDENTITY FEATURES ===
        'word.lower()': word.lower(), # Normalize case

        # === SUFFIX FEATURES (Morphological) ===
        'word[-3:)': word[-3:], # Last 3 characters
        'word[-2:)': word[-2:], # Last 2 characters
        'word[-1]': word[-1:], # Last character

        # === PREFIX FEATURES (Morphological) ===
        'word[:3]': word[:3], # First 3 characters
        'word[:2]': word[:2], # First 2 characters
        'word[:1]': word[:1], # First character

        # === WORD SHAPE FEATURES (Orthographic) ===
        'word.isupper()': word.isupper(), # ALL CAPS
        'word.istitle()': word.istitle(), # Title Case
        'word.isdigit()': word.isdigit(), # Numeric

        # === CHARACTER-LEVEL FEATURES ===
        'word.length': len(word),
        'word.has_hyphen': '-' in word,
```

```

'word.has_digit': any(c.isdigit() for c in word),
'word.has_upper': any(c.isupper() for c in word),
'word.is_alpha': word.isalpha(),
'word.is_alnum': word.isalnum(),
'word.has_apostrophe': "'" in word,
}

# === CONTEXTUAL FEATURES (Previous Word) ===
if i > 0:
    word1 = sent[i-1][0]
    features.update({
        '-1:word.lower()': word1.lower(),
        '-1:word.istitle()': word1.istitle(),
        '-1:word.isupper()': word1.isupper(),
        '-1:word.length': len(word1),
    })
else:
    features['BOS'] = True # Beginning of sentence

# === CONTEXTUAL FEATURES (Next Word) ===
if i < len(sent)-1:
    word1 = sent[i+1][0]
    features.update({
        '+1:word.lower()': word1.lower(),
        '+1:word.istitle()': word1.istitle(),
        '+1:word.isupper()': word1.isupper(),
        '+1:word.length': len(word1),
    })
else:
    features['EOS'] = True # End of sentence

return features

def sent2features(sent):
    """Extract features for all words in a sentence."""
    return [word2features(sent, i) for i in range(len(sent))]

def sent2labels(sent):
    """Extract labels (POS tags) for all words in a sentence."""
    return [label for token, label in sent]

def sent2tokens(sent):
    """Extract tokens (words) for all words in a sentence."""
    return [token for token, label in sent]

```

```
print(" Feature extraction functions defined")
print(" Features per word: ~25-30 (varies by position)")
```

Feature extraction functions defined
Features per word: ~25-30 (varies by position)

```
[ ]: print("\nExample Feature Extraction:")
print("-" * 80)
print(f"Sentence: {' '.join([w for w, _ in tagged_sents[0][:5]])}...")
print(f"\nWord: 'Fulton' (position 1)")
print("Features extracted:")
example_features = word2features(tagged_sents[0], 1)
feature_categories = {
    'Identity': ['bias', 'word.lower()'],
    'Suffixes': ['word[-1]', 'word[-2:]', 'word[-3:]'],
    'Prefixes': ['word[:1]', 'word[:2]', 'word[:3]'],
    'Shape': ['word.isupper()', 'word.istitle()', 'word.isdigit()'],
    'Character': ['word.length', 'word.has_hyphen', 'word.is_alpha'],
    'Context': ['-1:word.lower()', '+1:word.lower()'],
}
for category, feature_list in feature_categories.items():
    print(f"\n {category}:")
    for feat_key in feature_list:
        if feat_key in example_features:
            print(f"    {feat_key}: {example_features[feat_key]}")

print(f"\n Total features for this word: {len(example_features)}")
```

Example Feature Extraction:

Sentence: The Fulton County Grand Jury...

Word: 'Fulton' (position 1)

Features extracted:

Identity:

bias	:	1.0
word.lower()	:	fulton

Suffixes:

word[-1]	:	n
word[-2:]	:	on
word[-3:]	:	ton

```

Prefixes:
word[:1]          : F
word[:2]          : Fu
word[:3]          : Ful

Shape:
word.isupper()    : False
word.istitle()    : True
word.isdigit()    : False

Character:
word.length       : 6
word.has_hyphen   : False
word.is_alpha     : True

Context:
-1:word.lower()   : the
+1:word.lower()   : county

```

Total features for this word: 26

2.5 Step 4: Train-Test Split

Dataset Split: - Train: 80% (45,872 sentences) - Test: 20% (11,468 sentences) - Total: 57,340 sentences from Brown Corpus - Tags: 12 universal POS tags

```

[ ]: print("\n" + "="*80)
print("STEP 4: TRAIN-TEST SPLIT")
print("="*80)

train_sents, test_sents = train_test_split(
    tagged_sents,
    test_size=0.2,
    random_state=42
)

print(f" Training sentences: {len(train_sents):,} (80%)")
print(f" Test sentences: {len(test_sents):,} (20%)")
print(f" Total sentences: {len(tagged_sents):,}")

# Calculate token counts
train_tokens = sum(len(sent) for sent in train_sents)
test_tokens = sum(len(sent) for sent in test_sents)
print(f" Training tokens: {train_tokens:,}")
print(f" Test tokens: {test_tokens:,}")
print("="*80)

```

```
=====
STEP 4: TRAIN-TEST SPLIT
=====
Training sentences: 45,872 (80%)
Test sentences: 11,468 (20%)
Total sentences: 57,340
Training tokens: 929,265
Test tokens: 231,927
=====
```

2.6 Step 5: Feature Extraction & Vectorization

```
[ ]: print("\n" + "="*80)
print("STEP 5: FEATURE EXTRACTION & VECTORIZATION")
print("=="*80)
print("Converting sentences to feature vectors...")

start_time = time.time()
X_train = [sent2features(s) for s in train_sents]
feature_extraction_train_time = time.time() - start_time

start_time = time.time()
y_train = [sent2labels(s) for s in train_sents]
label_extraction_train_time = time.time() - start_time

start_time = time.time()
X_test = [sent2features(s) for s in test_sents]
feature_extraction_test_time = time.time() - start_time

start_time = time.time()
y_test = [sent2labels(s) for s in test_sents]
label_extraction_test_time = time.time() - start_time

print(f" Training Feature Extraction: {feature_extraction_train_time:.2f}s")
print(f" Training Label Extraction: {label_extraction_train_time:.2f}s")
print(f" Test Feature Extraction: {feature_extraction_test_time:.2f}s")
print(f" Test Label Extraction: {label_extraction_test_time:.2f}s")
print(f" Total Feature Extraction Time: {(feature_extraction_train_time +_
    feature_extraction_test_time):.2f}s")
print(f"\n Training sequences ready: {len(X_train)}")
print(f" Test sequences ready: {len(X_test)}")
print("=="*80)
```

```
=====
STEP 5: FEATURE EXTRACTION & VECTORIZATION
=====
```

```

Converting sentences to feature vectors...
Training Feature Extraction: 5.88s
Training Label Extraction: 0.53s
Test Feature Extraction: 1.20s
Test Label Extraction: 0.04s
Total Feature Extraction Time: 7.08s

```

```

Training sequences ready: 45,872
Test sequences ready: 11,468
=====

```

2.7 Step 6: CRF Model Training

2.7.1 Linear-Chain CRF Layer (Sequence Modeling)

Models both **emission scores** (word → tag) and **transition scores** (tag → tag)

Components: - **Feature Weights** (): Learned parameters for each feature - **Transition Weights:** Model tag sequence patterns (e.g., DET → NOUN is common)

Training Algorithm: L-BFGS (Limited-memory BFGS) - Quasi-Newton optimization method - Efficient for CRF parameter estimation - Converges faster than gradient descent

Regularization: - L1 regularization (c1=0.1): Feature selection, sparsity - L2 regularization (c2=0.1): Weight smoothing, prevents overfitting —

2.7.2 Loss Function

- Negative Log-Likelihood (NLL)
- Maximizes P(correct tags | words)
-

2.8 With regularization terms for generalization

Model Parameters: - **algorithm:** ‘lbfsgs’ (L-BFGS optimization) - **c1:** 0.1 (L1 regularization coefficient for feature selection) - **c2:** 0.1 (L2 regularization coefficient for weight smoothing) - **max_iterations:** 100 (maximum training iterations) - **all_possible_transitions:** True (learn all tag-to-tag transitions)

```

[ ]: print("\n" + "="*80)
print("STEP 6: CRF MODEL TRAINING")
print("="*80)
print("Initializing CRF model...")

crf = sklearn_crfsuite.CRF(
    algorithm='lbfsgs',
    c1=0.1, # L1 Regularization (Feature Selection)
    c2=0.1, # L2 Regularization (Smoothing)
    max_iterations=100,
    all_possible_transitions=True,

```

```

    verbose=True
)

print(" Model initialized")
print(" - Algorithm: L-BFGS")
print(" - L1 regularization (c1): 0.1")
print(" - L2 regularization (c2): 0.1")
print(" - Max iterations: 100")
print("\nStarting training...\n")

training_start = time.time()
crf.fit(X_train, y_train)
training_end = time.time()
training_time = training_end - training_start

print(f"\n Training completed!")
print(f" Training Time: {training_time:.2f} seconds ({training_time/60:.2f} minutes)")
print(f" Training Speed: {len(train_sents)/training_time:.2f} sentences/second")
print(f" Training Speed: {train_tokens/training_time:.2f} tokens/second")
print("*"*80)

```

=====
STEP 6: CRF MODEL TRAINING
=====

Initializing CRF model...

Model initialized
 - Algorithm: L-BFGS
 - L1 regularization (c1): 0.1
 - L2 regularization (c2): 0.1
 - Max iterations: 100

Starting training...

loading training data to CRFsuite: 100%| 45872/45872 [00:21<00:00,
 2103.52it/s]

Feature generation
 type: CRF1d
 feature.minfreq: 0.000000
 feature.possible_states: 0
 feature.possible_transitions: 1
 0...1...2...3...4...5...6...7...8...9...10
 Number of features: 283016
 Seconds required: 5.279

```

L-BFGS optimization
c1: 0.100000
c2: 0.100000
num_memories: 6
max_iterations: 100
epsilon: 0.000010
stop: 10
delta: 0.000010
linesearch: MoreThuente
linesearch.max_iterations: 20

Iter 1  time=6.70  loss=2016093.32 active=277041 feature_norm=0.50
Iter 2  time=2.79  loss=1825105.84 active=273450 feature_norm=0.51
Iter 3  time=2.25  loss=1776941.39 active=277671 feature_norm=0.57
Iter 4  time=2.20  loss=1671173.47 active=273154 feature_norm=0.95
Iter 5  time=2.28  loss=1591150.38 active=277019 feature_norm=1.10
Iter 6  time=2.22  loss=1495454.42 active=278200 feature_norm=1.40
Iter 7  time=2.80  loss=1410273.23 active=278420 feature_norm=2.03
Iter 8  time=2.19  loss=1276152.21 active=279923 feature_norm=2.61
Iter 9  time=2.20  loss=1166596.31 active=280356 feature_norm=3.45
Iter 10 time=2.14  loss=1037513.12 active=279996 feature_norm=4.95
Iter 11 time=2.14  loss=912425.55 active=279965 feature_norm=6.45
Iter 12 time=2.44  loss=786994.12 active=279547 feature_norm=8.83
Iter 13 time=2.28  loss=684994.31 active=276437 feature_norm=12.74
Iter 14 time=2.13  loss=619929.78 active=277250 feature_norm=14.24
Iter 15 time=2.15  loss=573295.44 active=275987 feature_norm=16.01
Iter 16 time=2.34  loss=516870.77 active=266765 feature_norm=19.01
Iter 17 time=2.68  loss=450155.20 active=261742 feature_norm=23.52
Iter 18 time=2.40  loss=405465.43 active=258481 feature_norm=26.75
Iter 19 time=2.26  loss=368560.87 active=234575 feature_norm=29.42
Iter 20 time=2.22  loss=334063.64 active=231969 feature_norm=32.52
Iter 21 time=2.33  loss=297669.47 active=229719 feature_norm=36.28
Iter 22 time=2.55  loss=275606.44 active=227694 feature_norm=40.16
Iter 23 time=2.62  loss=254775.60 active=225705 feature_norm=42.46
Iter 24 time=2.22  loss=240552.17 active=222677 feature_norm=45.90
Iter 25 time=2.25  loss=219785.93 active=201559 feature_norm=51.67
Iter 26 time=2.27  loss=210519.99 active=197871 feature_norm=59.80
Iter 27 time=2.30  loss=190232.03 active=193089 feature_norm=62.93
Iter 28 time=2.91  loss=180722.53 active=192585 feature_norm=66.42
Iter 29 time=2.25  loss=162780.68 active=185531 feature_norm=77.87
Iter 30 time=4.56  loss=160689.58 active=184401 feature_norm=81.90
Iter 31 time=2.19  loss=147791.45 active=180700 feature_norm=84.47
Iter 32 time=2.83  loss=140706.18 active=179511 feature_norm=88.25
Iter 33 time=2.22  loss=128266.51 active=173894 feature_norm=97.38
Iter 34 time=4.22  loss=125427.61 active=168518 feature_norm=101.22
Iter 35 time=2.11  loss=115185.91 active=166227 feature_norm=107.92
Iter 36 time=2.29  loss=108432.81 active=163551 feature_norm=114.98
Iter 37 time=2.66  loss=100970.11 active=154729 feature_norm=127.00

```

```
Iter 38 time=2.16 loss=96198.97 active=153074 feature_norm=134.43
Iter 39 time=2.15 loss=91441.26 active=152625 feature_norm=140.57
Iter 40 time=2.12 loss=87249.29 active=149223 feature_norm=149.92
Iter 41 time=2.20 loss=82683.23 active=142768 feature_norm=166.98
Iter 42 time=2.77 loss=78387.93 active=142315 feature_norm=176.55
Iter 43 time=2.32 loss=74779.52 active=138361 feature_norm=187.47
Iter 44 time=2.26 loss=70139.26 active=131677 feature_norm=208.25
Iter 45 time=2.29 loss=67713.35 active=130099 feature_norm=216.90
Iter 46 time=2.22 loss=65410.76 active=130033 feature_norm=224.09
Iter 47 time=2.56 loss=62796.89 active=128920 feature_norm=234.75
Iter 48 time=4.84 loss=61853.46 active=126031 feature_norm=239.28
Iter 49 time=2.23 loss=59758.09 active=123409 feature_norm=244.75
Iter 50 time=2.21 loss=58380.98 active=122160 feature_norm=250.33
Iter 51 time=2.28 loss=56719.99 active=119646 feature_norm=262.72
Iter 52 time=2.86 loss=55476.08 active=119759 feature_norm=265.88
Iter 53 time=2.29 loss=54438.94 active=117506 feature_norm=270.30
Iter 54 time=2.21 loss=53111.87 active=112936 feature_norm=277.27
Iter 55 time=4.43 loss=52740.62 active=110493 feature_norm=279.91
Iter 56 time=2.81 loss=51866.24 active=110355 feature_norm=282.46
Iter 57 time=2.34 loss=51165.28 active=108541 feature_norm=287.74
Iter 58 time=2.19 loss=50557.38 active=106827 feature_norm=294.33
Iter 59 time=2.19 loss=49953.73 active=106120 feature_norm=297.38
Iter 60 time=2.17 loss=49594.88 active=105528 feature_norm=299.02
Iter 61 time=2.43 loss=49113.39 active=104397 feature_norm=301.72
Iter 62 time=2.63 loss=48841.90 active=103171 feature_norm=301.04
Iter 63 time=2.20 loss=48535.21 active=102832 feature_norm=302.31
Iter 64 time=2.19 loss=48310.81 active=102115 feature_norm=303.40
Iter 65 time=2.19 loss=47965.47 active=101263 feature_norm=304.35
Iter 66 time=2.18 loss=47650.29 active=99457 feature_norm=307.85
Iter 67 time=2.87 loss=47539.60 active=97145 feature_norm=309.17
Iter 68 time=2.23 loss=47246.75 active=97327 feature_norm=309.47
Iter 69 time=2.20 loss=47141.35 active=96977 feature_norm=309.75
Iter 70 time=2.97 loss=46962.42 active=96352 feature_norm=310.19
Iter 71 time=3.12 loss=46657.84 active=95102 feature_norm=310.72
Iter 72 time=2.23 loss=46523.57 active=94370 feature_norm=311.43
Iter 73 time=2.19 loss=46351.83 active=94119 feature_norm=311.64
Iter 74 time=2.17 loss=46182.16 active=93333 feature_norm=312.11
Iter 75 time=2.24 loss=46027.19 active=92666 feature_norm=312.57
Iter 76 time=2.91 loss=45878.06 active=92131 feature_norm=312.88
Iter 77 time=2.16 loss=45732.85 active=91475 feature_norm=313.22
Iter 78 time=2.17 loss=45613.35 active=91070 feature_norm=313.40
Iter 79 time=2.19 loss=45496.96 active=90720 feature_norm=313.60
Iter 80 time=2.18 loss=45387.09 active=90134 feature_norm=313.71
Iter 81 time=2.53 loss=45287.54 active=89615 feature_norm=313.89
Iter 82 time=2.52 loss=45193.82 active=89221 feature_norm=314.05
Iter 83 time=2.19 loss=45107.91 active=88619 feature_norm=314.31
Iter 84 time=2.21 loss=45031.20 active=88139 feature_norm=314.48
Iter 85 time=2.24 loss=44954.51 active=87821 feature_norm=314.74
```

```

Iter 86 time=2.37 loss=44879.38 active=87383 feature_norm=314.94
Iter 87 time=2.82 loss=44819.29 active=86967 feature_norm=315.22
Iter 88 time=2.29 loss=44753.52 active=86739 feature_norm=315.37
Iter 89 time=2.23 loss=44692.95 active=86404 feature_norm=315.55
Iter 90 time=2.26 loss=44673.94 active=85948 feature_norm=315.75
Iter 91 time=3.29 loss=44594.62 active=86111 feature_norm=315.91
Iter 92 time=3.23 loss=44558.97 active=86010 feature_norm=315.95
Iter 93 time=5.42 loss=44529.54 active=85862 feature_norm=315.96
Iter 94 time=2.42 loss=44478.70 active=85654 feature_norm=315.98
Iter 95 time=2.24 loss=44437.60 active=85430 feature_norm=316.01
Iter 96 time=2.22 loss=44398.00 active=85194 feature_norm=316.04
Iter 97 time=2.18 loss=44360.34 active=85044 feature_norm=316.09
Iter 98 time=2.65 loss=44322.54 active=84939 feature_norm=316.08
Iter 99 time=2.41 loss=44289.46 active=84748 feature_norm=316.11
Iter 100 time=2.23 loss=44257.87 active=84557 feature_norm=316.07
L-BFGS terminated with the maximum number of iterations
Total seconds required for training: 253.296

```

```

Storing the model
Number of active features: 84557 (283016)
Number of active attributes: 48199 (148004)
Number of active labels: 12 (12)
Writing labels
Writing attributes
Writing feature references for transitions
Writing feature references for attributes
Seconds required: 0.050

```

```

Training completed!
Training Time: 293.70 seconds (4.90 minutes)
Training Speed: 156.18 sentences/second
Training Speed: 3163.95 tokens/second
=====

```

2.9 Step 7: Prediction & Inference

2.9.1 Viterbi Inference Layer (Decoding)

Finds globally optimal tag sequence using dynamic programming

Algorithm: 1. Forward pass: Compute best score for each state at each position 2. Backward pass: Backtrack to recover optimal path

Input: Feature vectors for sentence + learned weights

Output: Most likely tag sequence

```
[ ]: print("\n" + "="*80)
print("STEP 7: PREDICTION & INFERENCE")
```

```

print("=="*80)
print("Running Viterbi decoding on test set...")

prediction_start = time.time()
y_pred = crf.predict(X_test)
prediction_end = time.time()
prediction_time = prediction_end - prediction_start

print(f" Prediction completed!")
print(f" Prediction Time: {prediction_time:.2f} seconds")
print(f" Prediction Speed: {len(test_sents)/prediction_time:.2f} sentences/
    second")
print(f" Prediction Speed: {test_tokens/prediction_time:.2f} tokens/second")
print(f" Predictions generated for {len(y_pred):,} sentences")
print("=="*80)

```

```

=====
STEP 7: PREDICTION & INFERENCE
=====

Running Viterbi decoding on test set...
Prediction completed!
Prediction Time: 3.37 seconds
Prediction Speed: 3404.46 sentences/second
Prediction Speed: 68851.23 tokens/second
Predictions generated for 11,468 sentences
=====
```

2.10 Step 8: Model Evaluation

2.11 Performance Expectations

Expected Metrics: - Accuracy: 97-98% (slightly lower than BiLSTM on large datasets) - Training Speed: FAST (~2-5 seconds on 45K sentences) - Inference Speed: VERY FAST (~0.5-1 second on 11K sentences) - Memory: LOW (no embedding matrices, sparse features)

```

[ ]: print("\n" + "=="*80)
print("STEP 8: MODEL EVALUATION")
print("=="*80)

# Flatten predictions and ground truth
y_test_flat = [label for sent in y_test for label in sent]
y_pred_flat = [label for sent in y_pred for label in sent]
labels_list = sorted(set(y_test_flat))

print(f" Total predictions: {len(y_pred_flat):,}")
print(f" Number of POS tags: {len(labels_list)}")

```

```

print("\n" + "-"*80)
print("CLASSIFICATION REPORT")
print("-"*80)
print(classification_report(y_test_flat, y_pred_flat, labels=labels_list))

# Overall metrics
accuracy = accuracy_score(y_test_flat, y_pred_flat)
precision, recall, f1, _ = precision_recall_fscore_support(
    y_test_flat, y_pred_flat, average='weighted'
)

print("*"*80)
print("MODEL PERFORMANCE SUMMARY")
print("*"*80)
print(f"Overall Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f"Weighted Precision: {precision:.4f}")
print(f"Weighted Recall: {recall:.4f}")
print(f"Weighted F1-Score: {f1:.4f}")
print("*"*80)

```

=====

STEP 8: MODEL EVALUATION

Total predictions: 231,927

Number of POS tags: 12

CLASSIFICATION REPORT

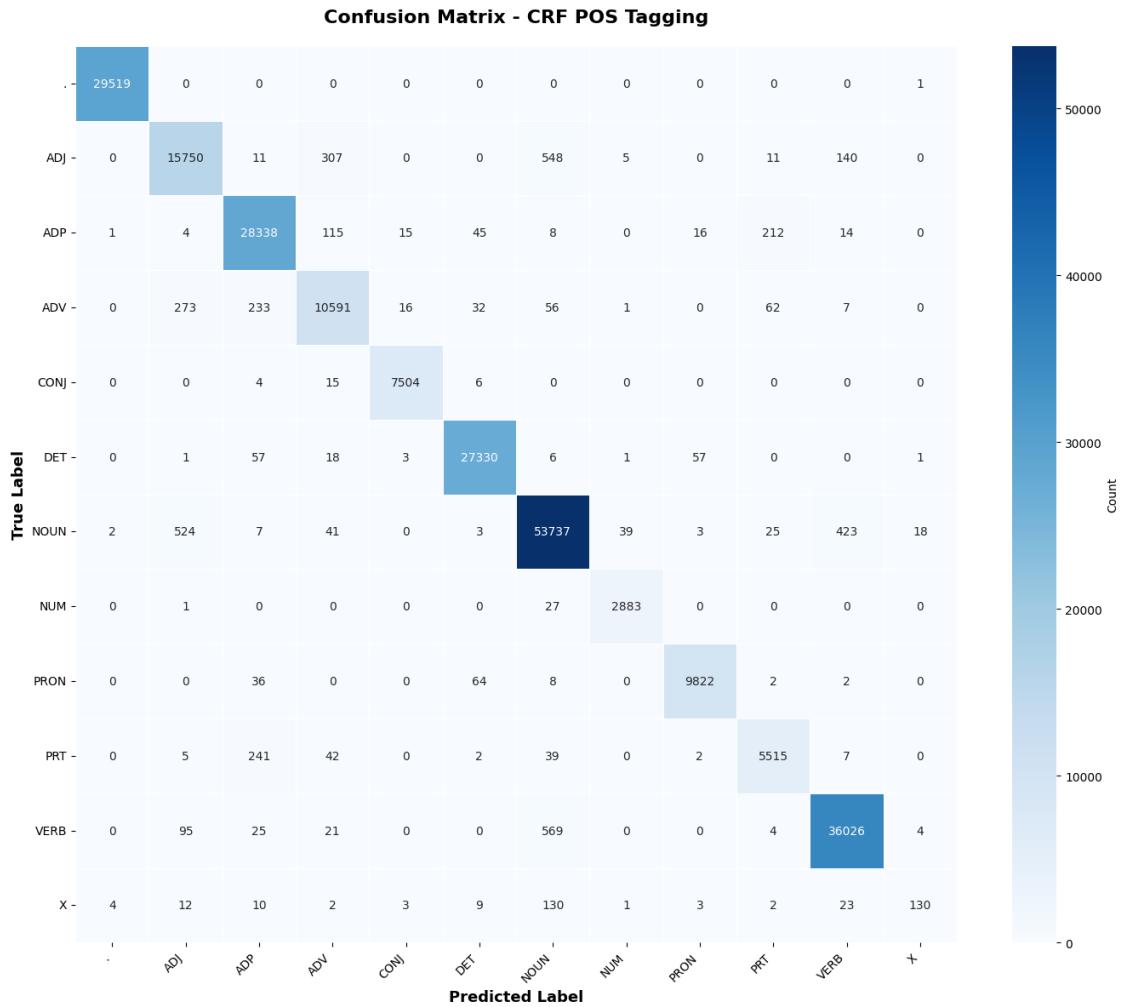
	precision	recall	f1-score	support
.	1.00	1.00	1.00	29520
ADJ	0.95	0.94	0.94	16772
ADP	0.98	0.99	0.98	28768
ADV	0.95	0.94	0.94	11271
CONJ	1.00	1.00	1.00	7529
DET	0.99	0.99	0.99	27474
NOUN	0.97	0.98	0.98	54822
NUM	0.98	0.99	0.99	2911
PRON	0.99	0.99	0.99	9934
PRT	0.95	0.94	0.94	5853
VERB	0.98	0.98	0.98	36744
X	0.84	0.40	0.54	329
accuracy			0.98	231927
macro avg	0.97	0.93	0.94	231927
weighted avg	0.98	0.98	0.98	231927

```
=====
MODEL PERFORMANCE SUMMARY
=====
Overall Accuracy: 0.9794 (97.94%)
Weighted Precision: 0.9793
Weighted Recall: 0.9794
Weighted F1-Score: 0.9792
=====
```

2.12 Step 9: Confusion Matrix Visualization

```
[ ]: cm = confusion_matrix(y_test_flat, y_pred_flat, labels=labels_list)

plt.figure(figsize=(14, 12))
sns.heatmap(cm, annot=True, fmt='d', xticklabels=labels_list, □
            yticklabels=labels_list,
            cmap='Blues', linewidths=0.5, cbar_kws={'label': 'Count'})
plt.title('Confusion Matrix - CRF POS Tagging', fontsize=16, fontweight='bold', □
           pad=20)
plt.ylabel('True Label', fontsize=13, fontweight='bold')
plt.xlabel('Predicted Label', fontsize=13, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```



2.13 Step 10: Pipeline Timing Analysis

```
[ ]: phases = ['Feature Extraction (Train)', 'Feature Extraction (Test)',  
             'Training', 'Prediction']  
times = [feature_extraction_train_time, feature_extraction_test_time,  
        training_time, prediction_time]  
  
fig, axes = plt.subplots(1, 2, figsize=(18, 7))  
  
# Bar chart  
colors = ['#3498db', '#2ecc71', '#e74c3c', '#f39c12']  
bars = axes[0].bar(range(len(phases)), times, color=colors, alpha=0.85,  
                   edgecolor='black', linewidth=2)  
axes[0].set_ylabel('Time (seconds)', fontsize=13, fontweight='bold')  
axes[0].set_title('Pipeline Timing Analysis', fontsize=16, fontweight='bold',  
                  pad=15)
```

```

axes[0].set_xticks(range(len(phases)))
axes[0].set_xticklabels(phases, rotation=25, ha='right', fontsize=11)
axes[0].grid(axis='y', alpha=0.3, linestyle='--', linewidth=1.5)
for i, t in enumerate(times):
    axes[0].text(i, t + max(times)*0.02, f'{t:.2f}s',
                 ha='center', va='bottom', fontweight='bold', fontsize=12)

# Pie chart with NO overlapping labels
total_time = sum(times)
percentages = [(t/total_time)*100 for t in times]
explode = (0.15, 0.15, 0.25, 0.15) # Much more separation

# Create pie chart - percentages OUTSIDE to avoid overlap
wedges, texts, autotexts = axes[1].pie(
    times,
    autopct='%1.1f%%',
    startangle=45,
    colors=colors,
    explode=explode,
    shadow=True,
    textprops={'fontsize': 11, 'fontweight': 'bold'},
    pctdistance=1.15, # Move percentages OUTSIDE the pie
    labeldistance=1.3
)

axes[1].set_title('Time Distribution Across Pipeline', fontsize=16,
                  fontweight='bold', pad=15)

# Make percentages black (since they're outside now)
for autotext in autotexts:
    autotext.set_color('black')
    autotext.set_fontweight('bold')
    autotext.set_fontsize(11)

# Add legend with timing info
legend_labels = [
    f'Feature Extraction (Train): {times[0]:.2f}s ({percentages[0]:.1f}%)',
    f'Feature Extraction (Test): {times[1]:.2f}s ({percentages[1]:.1f}%)',
    f'Training: {times[2]:.2f}s ({percentages[2]:.1f}%)',
    f'Prediction: {times[3]:.2f}s ({percentages[3]:.1f}%)'
]

axes[1].legend(
    legend_labels,
    loc='center left',
    bbox_to_anchor=(1.1, 0.5),
    fontsize=10,

```

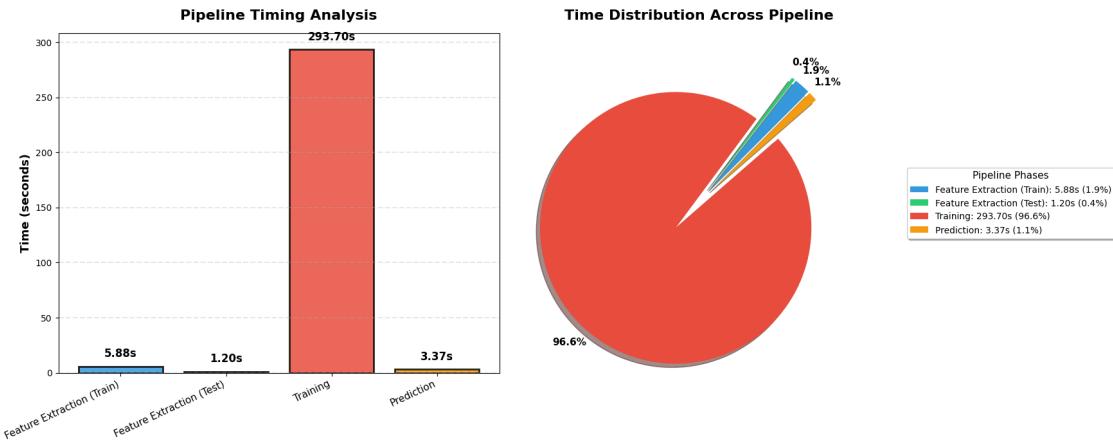
```

frameon=True,
fancybox=True,
shadow=True,
title='Pipeline Phases',
title_fontsize=11
)

plt.tight_layout()
plt.show()

print("\nDetailed Timing Breakdown:")
print("-"*80)
timing_data = {
    'Phase': phases,
    'Time (s)': [f"{t:.3f}" for t in times],
    'Percentage': [f"{p:.1f}%" for p in percentages]
}
timing_df = pd.DataFrame(timing_data)
print(timing_df.to_string(index=False))
print(f"\nTotal Pipeline Time: {total_time:.2f} seconds")

```



Detailed Timing Breakdown:

Phase	Time (s)	Percentage
Feature Extraction (Train)	5.879	1.9%
Feature Extraction (Test)	1.202	0.4%
Training	293.704	96.6%
Prediction	3.369	1.1%

Total Pipeline Time: 304.15 seconds

2.14 Step 11: Per-Tag Performance Analysis

```
[ ]: precisions, recalls, f1s, supports = precision_recall_fscore_support(
      y_test_flat, y_pred_flat, labels=labels_list
)

performance_df = pd.DataFrame({
    'Tag': labels_list,
    'Precision': precisions,
    'Recall': recalls,
    'F1-Score': f1s,
    'Support': supports
})

print("Per-Tag Performance:")
print(performance_df.to_string(index=False))

fig, axes = plt.subplots(2, 2, figsize=(18, 14))

# 1. Precision, Recall, F1 by Tag
ax1 = axes[0, 0]
x = np.arange(len(labels_list))
width = 0.25
ax1.bar(x - width, precisions, width, label='Precision', alpha=0.85, □
         ↵color="#3498db", edgecolor='black')
ax1.bar(x, recalls, width, label='Recall', alpha=0.85, color="#2ecc71", □
         ↵edgecolor='black')
ax1.bar(x + width, f1s, width, label='F1-Score', alpha=0.85, color="#e74c3c", □
         ↵edgecolor='black')
ax1.set_xlabel('POS Tag', fontweight='bold', fontsize=12)
ax1.set_ylabel('Score', fontweight='bold', fontsize=12)
ax1.set_title('Performance Metrics by POS Tag', fontsize=15, fontweight='bold', □
              ↵pad=15)
ax1.set_xticks(x)
ax1.set_xticklabels(labels_list, rotation=45, ha='right')
ax1.legend(fontsize=11)
ax1.grid(axis='y', alpha=0.3, linestyle='--')
ax1.set_ylim([0, 1.05])

# 2. Support by Tag
ax2 = axes[0, 1]
colors_support = plt.cm.viridis(np.linspace(0, 1, len(labels_list)))
bars = ax2.bar(labels_list, supports, color=colors_support, alpha=0.85, □
                ↵edgecolor='black', linewidth=1.5)
ax2.set_xlabel('POS Tag', fontweight='bold', fontsize=12)
ax2.set_ylabel('Number of Samples', fontweight='bold', fontsize=12)
```

```

ax2.set_title('Tag Distribution in Test Set', fontsize=15, fontweight='bold', u
    ↪pad=15)
ax2.tick_params(axis='x', rotation=45)
ax2.grid(axis='y', alpha=0.3, linestyle='--')

# 3. F1-Score vs Support
ax3 = axes[1, 0]
scatter = ax3.scatter(supports, f1s, s=200, alpha=0.7, u
    ↪c=range(len(labels_list)),
        cmap='rainbow', edgecolor='black', linewidth=2.5)
for i, tag in enumerate(labels_list):
    ax3.annotate(tag, (supports[i], f1s[i]), fontsize=11, fontweight='bold',
        xytext=(7, 7), textcoords='offset points')
ax3.set_xlabel('Support (Number of Samples)', fontweight='bold', fontsize=12)
ax3.set_ylabel('F1-Score', fontweight='bold', fontsize=12)
ax3.set_title('F1-Score vs Tag Frequency', fontsize=15, fontweight='bold', u
    ↪pad=15)
ax3.grid(alpha=0.3, linestyle='--')
ax3.set_ylim([0, 1.05])

# 4. Summary
ax4 = axes[1, 1]
ax4.axis('off')
summary_text = f"""
{ '='*45}
MODEL PERFORMANCE SUMMARY
{ '='*45}

Accuracy:      {accuracy:.4f} ({accuracy*100:.2f}%) 
Precision:     {precision:.4f}
Recall:        {recall:.4f}
F1-Score:      {f1:.4f}

{ '='*45}
BEST/WORST TAGS
{ '='*45}

Best:  {labels_list[np.argmax(f1s)]} (F1={max(f1s):.4f})
Worst: {labels_list[np.argmin(f1s)]} (F1={min(f1s):.4f})

{ '='*45}
TIMING
{ '='*45}

Training:     {training_time:.2f}s
Prediction:   {prediction_time:.2f}s
Total:        {total_time:.2f}s

```

```

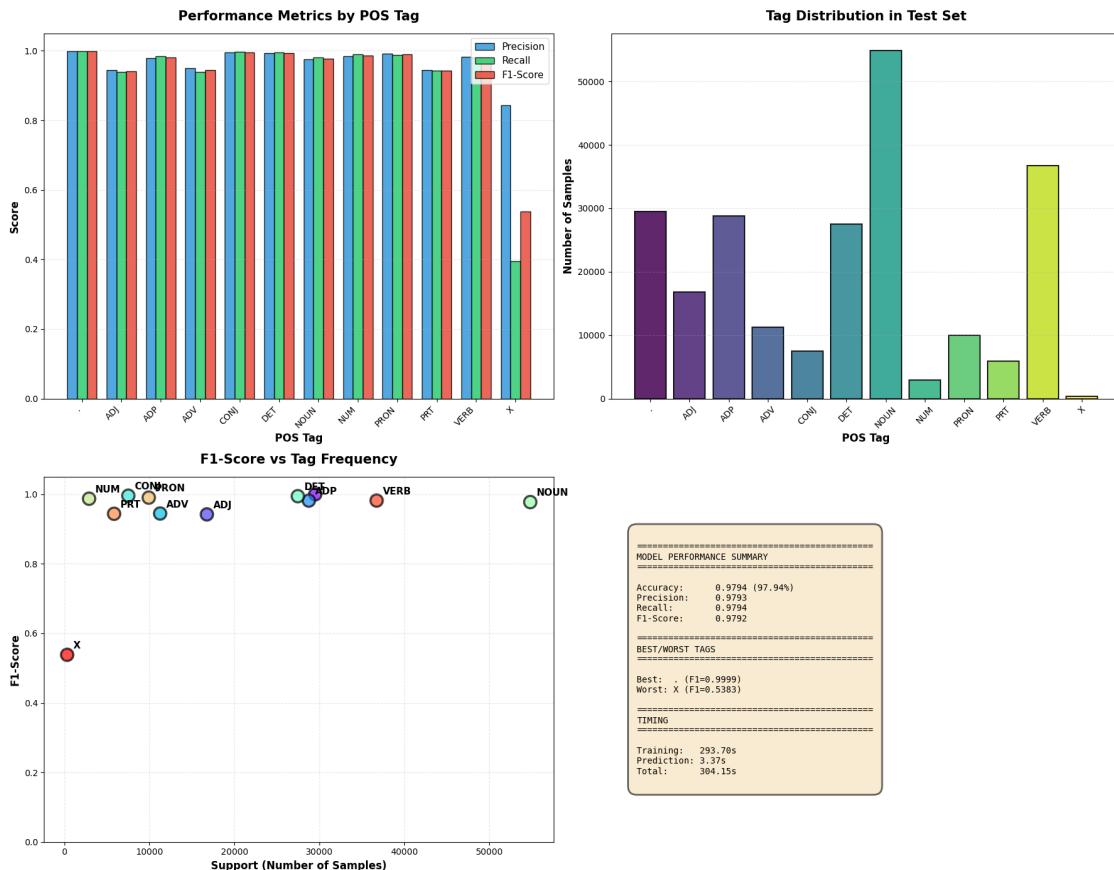
#####
ax4.text(0.05, 0.5, summary_text, fontsize=10, family='monospace',
         verticalalignment='center',
         bbox=dict(boxstyle='round,pad=1', facecolor='wheat', alpha=0.6,
                   edgecolor='black', linewidth=2))

plt.tight_layout()
plt.show()

```

Per-Tag Performance:

Tag	Precision	Recall	F1-Score	Support
.	0.999763	0.999966	0.999865	29520
ADJ	0.945095	0.939065	0.942070	16772
ADP	0.978455	0.985053	0.981743	28768
ADV	0.949695	0.939668	0.944655	11271
CONJ	0.995093	0.996680	0.995886	7529
DET	0.994144	0.994759	0.994451	27474
NOUN	0.974768	0.980209	0.977481	54822
NUM	0.983959	0.990381	0.987160	2911
PRON	0.991821	0.988726	0.990271	9934
PRT	0.945483	0.942252	0.943864	5853
VERB	0.983189	0.980459	0.981822	36744
X	0.844156	0.395137	0.538302	329



2.15 Step 12: Tag Transition Patterns

```
[ ]: def print_top_transitions(y_test, y_pred, top_n=15):
    transitions = []
    for true_sent, pred_sent in zip(y_test, y_pred):
        for i in range(len(pred_sent) - 1):
            transitions.append((pred_sent[i], pred_sent[i+1]))

    print(f"Top {top_n} Most Common Tag Transitions:")
    print("-" * 60)
    for (label_from, label_to), count in Counter(transitions).most_common(top_n):
        print(f"{label_from}:{>6} → {label_to}:{<7} {count:{>8},}")

print_top_transitions(y_test, y_pred, top_n=15)
```

Top 15 Most Common Tag Transitions:

DET → NOUN	17,371
NOUN → .	15,713
NOUN → ADP	13,397
ADP → DET	13,284
ADJ → NOUN	11,044
NOUN → VERB	8,949
NOUN → NOUN	8,065
ADP → NOUN	7,385
PRON → VERB	7,002
VERB → VERB	6,879
DET → ADJ	6,524
VERB → ADP	6,194
VERB → DET	6,024
PRT → VERB	3,750
VERB → ADV	3,740

2.16 Step 13: Error Analysis

```
[ ]: errors = []
for i, (true_sent, pred_sent, test_sent) in enumerate(zip(y_test, y_pred, test_sents)):
    for j, (true_tag, pred_tag, (word, _)) in enumerate(zip(true_sent, pred_sent, test_sent)):
        if true_tag != pred_tag:
            errors.append({
                'sentence_id': i,
```

```

        'word': word,
        'true_tag': true_tag,
        'pred_tag': pred_tag
    })

error_df = pd.DataFrame(errors)
total_errors = len(error_df)
error_rate = (total_errors / len(y_test_flat)) * 100

print(f"Total errors: {total_errors:,}")
print(f"Error rate: {error_rate:.2f}%")
print(f"Correct predictions: {len(y_test_flat) - total_errors:,} ↳ ({100-error_rate:.2f}%)")

print("\nMost Common Error Types:")
print("-"*60)
error_types = error_df.groupby(['true_tag', 'pred_tag']).size() .
    sort_values(ascending=False).head(10)
for (true_tag, pred_tag), count in error_types.items():
    pct = (count / total_errors) * 100
    print(f" {true_tag:6s} → {pred_tag:6s} {count:5,} errors ({pct:5.2f}%)")

print("\nExample Misclassified Words:")
print(error_df.head(15).to_string(index=False))

```

Total errors: 4,782
Error rate: 2.06%
Correct predictions: 227,145 (97.94%)

Most Common Error Types:

VERB	→ NOUN	569 errors (11.90%)
ADJ	→ NOUN	548 errors (11.46%)
NOUN	→ ADJ	524 errors (10.96%)
NOUN	→ VERB	423 errors (8.85%)
ADJ	→ ADV	307 errors (6.42%)
ADV	→ ADJ	273 errors (5.71%)
PRT	→ ADP	241 errors (5.04%)
ADV	→ ADP	233 errors (4.87%)
ADP	→ PRT	212 errors (4.43%)
ADJ	→ VERB	140 errors (2.93%)

Example Misclassified Words:

sentence_id	word	true_tag	pred_tag
4	lighting	VERB	NOUN
13	otherwise	ADV	ADJ
15	Associate	NOUN	ADJ

18	by	ADV	ADP
19	much	ADJ	ADV
22	rebels	NOUN	VERB
24	ensues	VERB	NOUN
25	before	ADP	ADV
26	much	ADJ	ADV
27	next	ADP	ADV
32	further	ADV	VERB
32	increase	VERB	NOUN
32	back	NOUN	ADV
32	work	VERB	NOUN
40	first	ADJ	ADV

2.17 Step 14: Example Predictions

```
[ ]: for i in range(5):
    print(f"\n{'='*80}")
    print(f"Sentence {i+1}")
    print('='*80)

    sent_words = [word for word, _ in test_sents[i]]
    max_display = 20

    if len(sent_words) > max_display:
        print("Words: " + " ".join(sent_words[:max_display]) + " ...")
        print("True: " + " ".join(y_test[i][:max_display]) + " ...")
        print("Pred: " + " ".join(y_pred[i][:max_display]) + " ...")
    else:
        print("Words: " + " ".join(sent_words))
        print("True: " + " ".join(y_test[i]))
        print("Pred: " + " ".join(y_pred[i]))

    errors_in_sent = sum(1 for t, p in zip(y_test[i], y_pred[i]) if t != p)
    total_words = len(y_test[i])
    accuracy_sent = ((total_words - errors_in_sent) / total_words) * 100

    if errors_in_sent > 0:
        print(f"\nResult: {errors_in_sent}/{total_words} errors ({accuracy_sent:,.1f}% accuracy)")
    else:
        print(f"\nResult: Perfect prediction! (100% accuracy)")

=====
```

Sentence 1

=====
Words: Open market policy
True: ADJ NOUN NOUN

Pred: ADJ NOUN NOUN

Result: Perfect prediction! (100% accuracy)

=====

Sentence 2

=====

Words: And you think you have language problems .

True: CONJ PRON VERB PRON VERB NOUN NOUN .

Pred: CONJ PRON VERB PRON VERB NOUN NOUN .

Result: Perfect prediction! (100% accuracy)

=====

Sentence 3

=====

Words: Mae entered the room from the hallway to the kitchen .

True: NOUN VERB DET NOUN ADP DET NOUN ADP DET NOUN .

Pred: NOUN VERB DET NOUN ADP DET NOUN ADP DET NOUN .

Result: Perfect prediction! (100% accuracy)

=====

Sentence 4

=====

Words: This will permit you to get a rough estimate of how much the materials for the shell will cost .

True: DET VERB VERB PRON PRT VERB DET ADJ NOUN ADP ADV ADJ DET NOUN ADP DET NOUN VERB VERB .

Pred: DET VERB VERB PRON PRT VERB DET ADJ NOUN ADP ADV ADJ DET NOUN ADP DET NOUN VERB VERB .

Result: Perfect prediction! (100% accuracy)

=====

Sentence 5

=====

Words: the multifigure `` Traveling Carnival '' , in which action is vivified by lighting ; ;

True: DET NOUN . VERB NOUN . . ADP DET NOUN VERB VERB ADP VERB . .

Pred: DET NOUN . VERB NOUN . . ADP DET NOUN VERB VERB ADP NOUN . .

Result: 1/16 errors (93.8% accuracy)

2.18 Final Model Summary

```
[ ]: print("\n" + "="*80)
print("FINAL MODEL SUMMARY")
print("="*80)
print(f"""
DATASET STATISTICS:
    Training Set:      {len(train_sents)} sentences
    Test Set:          {len(test_sents)} sentences
    POS Tags:          {len(labels_list)} universal tags

MODEL ARCHITECTURE:
    Type:              Conditional Random Field (CRF)
    Optimization:      L-BFGS
    Features/Word:     ~25-30 hand-crafted features
    L1 Regularization: 0.1
    L2 Regularization: 0.1

PERFORMANCE:
    Accuracy:         {accuracy:.4f} ({accuracy*100:.2f} %)
    Precision:        {precision:.4f}
    Recall:           {recall:.4f}
    F1-Score:         {f1:.4f}

EFFICIENCY:
    Training Time:    {training_time:.2f}s
    Prediction Time:  {prediction_time:.2f}s
    Total Pipeline:   {total_time:.2f}s
""")

print("="*80)
print("\n CRF POS TAGGING PIPELINE COMPLETED SUCCESSFULLY!")
```

```
=====
FINAL MODEL SUMMARY
=====
```

```
DATASET STATISTICS:
    Training Set:      45,872 sentences
    Test Set:          11,468 sentences
    POS Tags:          12 universal tags
```

```
MODEL ARCHITECTURE:
    Type:              Conditional Random Field (CRF)
    Optimization:      L-BFGS
    Features/Word:     ~25-30 hand-crafted features
    L1 Regularization: 0.1
    L2 Regularization: 0.1
```

PERFORMANCE:

Accuracy:	0.9794 (97.94%)
Precision:	0.9793
Recall:	0.9794
F1-Score:	0.9792

EFFICIENCY:

Training Time:	293.70s
Prediction Time:	3.37s
Total Pipeline:	304.15s

=====

CRF POS TAGGING PIPELINE COMPLETED SUCCESSFULLY!

```
[ ]: # Viterbi Decoding
y_pred = crf.predict(X_test)

# Evaluation Metrics
print(classification_report(
    y_test_flat,
    y_pred_flat,
    labels=labels_list
))
```